

UPPSALA UNIVERSITY



PARALLEL AND DISTRIBUTED PROGRAMMING

1TD070

Project - report

Authors:

Csongor HORVÁTH

May 3, 2023

The problem setting

The problem was to implement the Parallel Quicksort algorithm using C and MPI. The algorithm (Algorithm 1) and the usage of the code, which includes the 3 given parameter, the necessary outputs and the format of the I/O files was described in the assignment, so I won't describe them here.

Algorithm 1 The Parallel Quick-sort algorithm

1. Divide the data into p equal parts, one per process
2. Sort the data locally for each process
3. Perform global sort
 - 3.1 Select pivot element within each process set
 - 3.2 Locally in each process, divide the data into two sets according to the pivot (smaller or larger)
 - 3.3 split the processes into two groups and exchange data pairwise between them so that all processes in one group get data less than the pivot and the others get data larger than the pivot.
 - 3.4 Merge the two sets of numbers in each process into one sorted list
4. Repeat 3.1 - 3.4 recursively for each half until each group consists of one single process.

This concludes the brief description of the problem setting. Information about my implementation and further details can be found in the next section.

Implementation

For the implementations it could be assumed that p , so the number of processes is a power of 2. Now there are I/O processes, which similarly to the previous assignments is done by the root process. I use rank 0 as root. As this processes are very similar to the previous assignment I won't explain them in details. I use functions for the read and write process and the run times are collected using `MPI_Reduce` to get the value to print in the screen in the root process.

I wrote two function to check if p is a 2 power and if it is then get $\log_2(p)$, as these values needed in the project. Then I wrote a serial quicksort in 3 function which is needed in the 2nd step of the algorithm. And I implemented a merge function for the step 3.4 in the algorithm for testing. Also there is a `print_list` function which was used for testing.

Now let's talk about the assignment specific details. First note that even if the algorithm was given in a recursive form my implementation wasn't following a recursive way as knowing p the depth of the recursion is known to be $\log_2(p)$, so instead of doing things recursively I used a loop and not like in the serial implementation I also split the pivot element into one of the sub processes, so in the end it was enough to gather the local list from each processor in the root to get the final ordered list. My hope was that in this way the list doesn't build up in a recursive manner, therefore it could be faster even without the left of the pivot elements, which wouldn't count for much considering the large size of the problem comparison to the dept of the recursion.

One nice assumption would be that $p \mid n$, so n dividable by p , but as it was asked so my program is working for arbitrary n value in the following manner: if $n < p$ then as the number of cores are relatively small, so in this case I only do a serial sort in the root process. If $p \mid n$, then I simply use `MPI_Scatter` to distribute the list between the processes. In other cases I make a calculation and divide as equally as possible (max difference is 1 between the processes in the list size). I use `MPI_Scatterv` for it after calculating the necessary parameters.

After distribution comes the 2nd step of the algorithm, so the local serial sort. I used quicksort for this, which may not be the best choice. As it will be explained later this is the most crucial part of the whole process, so it

is recommended to choose a sorting algorithm here which is suitable for the use case of our application.

After this starts the recursive part, which is implemented in a loop as it's depth is given as I mention. The splitting of the groups are solved with the splitting of the communicator using `MPI_Comm_split` for splitting half each group from the previous state.

Now the next step is to get a pivot element, for this I use a function. It is a complicated process to get the pivot element as for this purpose the communication between the threads are required and also there are 3 case based on the user. Also the possibility that some of the processes has empty list in a given step makes this process more complicated as in this case the calculation of the pivot element is getting more complex with the consideration of which element is empty in the group.

The basic idea of the implementation is to gather all necessary data at one node in each groups and make the calculation locally there and finally distribute the pivot element in the group. Everything else is only the details for the lot of different cases and the simply calculation of the required values in each cases.

After the pivot element the splitting of the data comes as the following. First I calculate indexes, where the data smaller and larger than the pivot, and with which process should the informations be exchanged. Then I communicate the size of the exchange in both direction and then if the size is positive then I send and receive the data. I used `MPI_Isend` with `MPI_Wait` and `MPI_Recv` for communication here as the non-blocking send made the two way exchange more simple in code compared to the blocking, where I should pay attention to the order for the information exchange.

And finally I merge the list using a merge function. Note that the merge process is fast and simple as the two partial list are both ordered.

Now the above steps run in a loop $\log_2(p)$ times instead of the recursion until there is one process in each group, but the result is the same this way.

In the end I gather the length of the distributed lists and then gather the partial list from the processes in a way to get a fully ordered list from them

in the root process.

For the experiment note that in this case the time measurement includes the distribution and the gathering of the list.

Partitioning strategies

The following partitioning strategies are implemented in the project:

1. Select the median in one processor in each group of processors.
2. Select the median of all medians in each processor group.
3. Select the mean value of all medians in each processor group.

There are other possible pivot strategies, but as it will be explained in this algorithm it is not the most crucial part of the algorithm.

Performance experiment

First I would like to note that I didn't presented run times for the backward input files as in that case my choice of the serial sorting algorithm was pretty bad, since that is the worst case for the serial quick sort algorithm. Also note that even if for the given file the run time was too large, I made measurements for smaller backward files which was already slow with smaller number of integers.

Here I would like to mention that I made partial measurements and I find it that in most cases the majority of run times comes from the serial sorting. In the smaller (10^6) backwards case it was up to 95%. But even with the random large lists the serial sorting in the beginning mostly takes more than 50% of the run time. Therefore if we know something about our application it is crucial to choose a suitable sorting algorithm, or do something to avoid it's weaknesses. E.g: in this case if I would like to make it usable to the backward sorted lists, then I could shuffle the list elements before or during the distribution.

After discussing this I will start the discussion of my experiment which was made using the large input files (*input125000000.txt*, *input250000000.txt*, *input500000000.txt*, *input1000000000.txt*, *input2000000000.txt*).

pivot	-n 1	-n 2	-n 4	-n 8	-n 16	-n 32	-n 64	-n 128
1	14.8555	8.1279	4.6732	2.8762	1.6016	1.1451	2.5540	7.8167
2	14.8592	8.1570	4.6980	2.8975	1.6006	1.1614	3.0454	7.9721
3	14.8587	8.5494	5.7424	4.3978	3.4562	3.2196	4.6182	9.3997

Table 1: Time measurement for *input125000000.txt* in 2 node 32 core using pivot strategy 1,2,3

My first goal was to compare the 3 different pivot strategy. For this end I made measurements with all 3 pivot strategy with 2^i , $i = 0, 1, \dots, 7$ core. I used the smallest large input file. The run times can be seen in Table 1. For 64 and 128 core the measurements are not good as I only used 2 node (32 core) for measurements and these were run using `-oversubscribe` flag. But it is good to see that it is not worth to try using oversubscribe in this

application as there are projects, where it would further improve or at least not make the performance significantly worse as in this case.

And as for the pivot strategies the 1 and 2 seems to be very similar with 1 being a slightly, almost immeasurably faster, but for some reason 3 is significantly worse and the gap is getting bigger as using more core. Honestly I don't know why it is becoming worse, maybe it is slower to calculate, or it makes things worse in other way.

I want to note that again that the starting serial sort is the most time consuming in most time, as after that not considering the communication and the calculation of pivot I only merge sorted lists, which is a very fast thing to do even for large amount of number.

Strong scaling experiment

size 10^6	serial	-n 1	-n 2	-n 4	-n 8	-n 16	-n 32
125	14.5551	14.8526	8.1641	4.6673	2.8931	1.6104	1.1479
250	29.6503	30.4716	16.7777	9.6095	5.9846	3.2291	2.2321
500	61.6230	62.995683	34.8779	19.8823	12.3496	6.6020	4.4787
1000	132.3010	131.6470	70.9468	41.5660	25.3717	13.5538	9.0914
2000	280.9389	274.6057	148.3555	84.9485	53.5367	27.8613	18.5007

Table 2: Time measurement for strong scaling in 2 node 32 core using pivot strategy 1

For the strong scaling I made measurements with all of the large simple input files using 2^i , $i = 0, 1, \dots, 5$ cores. For reference I also made measurements with a simple serial implementation of the problem. As we can see the starting point in one core is not much worse then the specific serial implementation, it is even faster in some cases, which is probably a measurement error.

For better visualization I plotted the run times for most cases in below Figure 1 with the ideal speed up and the ratio between the ideal speed up and the real run time.

As it can be seen the strong scaling is relatively good, the ratio is similar as

I get in previous assignments and it is almost the same between the different files. The measurements is also similar in a way, that the ratio gets worse when reaching the limitation of our resources. My suggested reasons was explained for this in the previous assignment.

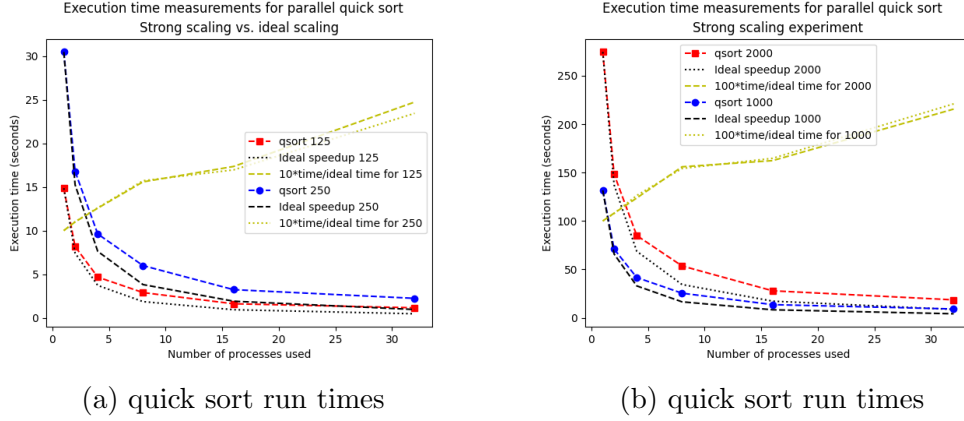


Figure 1: Strong scaling

Weak scaling experiment

core / size 10^6	$-n$ 1 / 125	$-n$ 2 / 250	$-n$ 4 / 500	$-n$ 8 / 1000	$-n$ 16 / 2000
quick sort	14.8526	16.7777	19.8823	25.3717	27.8613

Table 3: Time measurement for weak scaling in 2 node 32 core

For the weak scaling I used a subset of the run times measured above. The chosen run times can be seen in Table 3. Note that it is not easy to make a good weak scaling experiment as the scaling of the run time is not nicely describable. As in quick sort the worst case is $\mathcal{O}(n^2)$, but based on this we got a negative weak scaling as for this when changing input size to twice the size, then we should multiply the used number of cores with 4. A list of run times would be e.g:(125/1 : 14.8526, 250/4 : 9.6095 500/16 : 6.6020). So instead I used a linear scaling approach and I considered changing the file size to double as multiplying the used number of cores. Then I made a plot

in Figure 2. Here for ideal speed up I plotted the curve coming from the average run time $\mathcal{O}(n\log(n))$ and I used 1 as constant value in the \mathcal{O} run time.

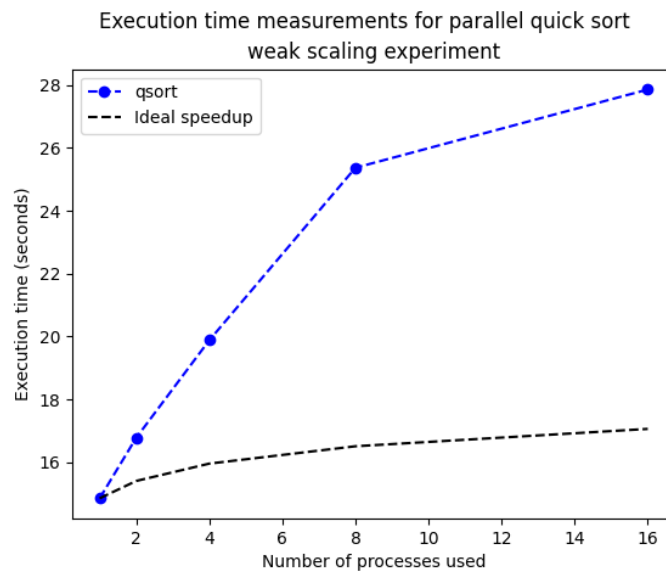


Figure 2: Weak scaling for quick sort

The weak scaling in this way looks quite bad, but I don't know for what I should measure the weak scaling I could use a squared scaling as that is the worst case, which would mean a line, but the real scaling is better then that. And I have no idea what constant I should use for the average case, so $\mathcal{O}(n\log(n))$ scaling. But I would say the tendency is a bit similar in the measured times with the $\mathcal{O}(n\log(n))$.

So in conclusion I would say that weak scaling experiment doesn't make much sense as we don't know what scaling we should expect from the algorithm for ideal case.

Summary

So in summary I would conclude that it can be useful to do the quick sort in a parallel manner as in experience I get a significant improvement until I get resource for it by using more core. And if in a really large scale application a sorting is needed, or in a relatively large scale it needed multiple times, then the parallel implementation can save us a lot of time.

I also would like to mention that a crucial part of this algorithm is the serial sort, so for real application we should be careful with it. Also the choice of the pivot can make impact on the run time.

There are also other possibilities to make a parallel implementation of quick sort. One example is to use the serial quick sort, but until there are empty processes let's start one of the recursive calls in new process.

Appendix

quicksort

```
1  #include <mpi.h>
2  #include <stdio.h>
3  #include <stdlib.h>
4  #include <math.h>
5
6
7  #define root 0
8  #define save 1
9
10 int read_input(const char *file_name, int **values) {
11     FILE *file;
12     if (NULL == (file = fopen(file_name, "r"))) {
13         perror("Couldn't open input file");
14         return -1;
15     }
16     int num_values;
17     if (EOF == fscanf(file, "%d", &num_values)) {
18         perror("Couldn't read element count from input file");
19         return -1;
20     }
21     if (NULL == (*values = malloc(num_values * sizeof(int)))) {
22         perror("Couldn't allocate memory for input");
23         return -1;
24     }
25     for (int i=0; i<num_values; i++) {
26         if (EOF == fscanf(file, "%d", &((*values)[i]))) {
27             perror("Couldn't read elements from input file");
28             return -1;
29         }
30     }
31     if (0 != fclose(file)) {
32         perror("Warning: couldn't close input file");
33     }
34     return num_values;
35 }
36
37 int write_output(char *file_name, const int *output, int num_values) {
38     FILE *file;
39     if (NULL == (file = fopen(file_name, "w"))) {
40         perror("Couldn't open output file");
41         return -1;
42     }
43     for (int i = 0; i < num_values; i++) {
44         if (0 > fprintf(file, "%d", output[i])) {
45             perror("Couldn't write to output file");
46         }
47     }
48     if (0 > fprintf(file, "\n")) {
49         perror("Couldn't write to output file");
50     }
51     if (0 != fclose(file)) {
52         perror("Warning: couldn't close output file");
53     }
54     return 0;
55 }
56
57 int is_power_of_two(int x){
58     return (x != 0) && ((x & (x - 1)) == 0);
59 }
60
61 int log_2(int x){
62     int count=0;
63     int pow=1;
64     while(x > pow){
65         count++;
66         pow*=2;
67     }
68     return count;
69 }
70 }
```

```

71
72 // Codes for serial sorting
73 int partition(int** arr, int low, int high) {
74     int pivot = *(*arr + high);
75     int i = (low - 1);
76
77     for (int j = low; j <= high - 1; j++) {
78         if (*(*arr + j) < pivot) {
79             i++;
80             int temp = *(*arr + i);
81             *(*arr + i) = *(*arr + j);
82             *(*arr + j) = temp;
83         }
84     }
85     int temp = *(*arr + i + 1);
86     *(*arr + i + 1) = *(*arr + high);
87     *(*arr + high) = temp;
88
89     return (i + 1);
90 }
91
92 void quickSort(int** arr, int low, int high) {
93     if (low < high) {
94         int pi = partition(arr, low, high);
95         quickSort(arr, low, pi - 1);
96         quickSort(arr, pi + 1, high);
97     }
98 }
99
100 void serial_sort(int** arr, int length) {
101     quickSort(arr, 0, length - 1);
102 }
103
104 int print_list(int* arr, int n){
105     for (size_t i = 0; i < n; i++)
106     {
107         printf("%d\\n", arr[i]);
108     }
109     printf("\\n");
110 }
111 }
112
113 int get_pivot(int median, int is_empty, MPIComm* comm, int rank, int size, int
    pivot_strategy){
114     int res; int all_empty; int sum;
115     int* medians=(int*)malloc(size*sizeof(int));
116     MPI_Allreduce(&is_empty, &all_empty, 1, MPI_INT, MPLSUM, *comm);
117     MPI_Gather(&median, 1, MPI_INT, medians, 1, MPI_INT, root, *comm);
118     if(!all_empty){
119         if(rank==root){
120             switch (pivot_strategy)
121             {
122                 case 1:
123                     res=median;
124                     break;
125
126                 case 2:
127                     serial_sort(&medians, size);
128                     res=medians[size/2];
129                     break;
130
131                 case 3:
132                     sum=0;
133                     for (int i = 0; i < size; i++)
134                         sum +=medians[i];
135                     res=sum/size;
136                     break;
137             }
138
139             for (int i = 0; i < size; i++)
140                 medians[i]=res;
141         }
142     }
143     else{ // case where some of the processes has currently empty sets
144         int* empties=(int*)malloc(size*sizeof(int));
145         MPI_Gather(&is_empty, 1, MPI_INT, empties, 1, MPI_INT, root, *comm);
146         if(rank==root){

```

```

147         res=0;
148         switch (pivot_strategy)
149         {
150             case 1:
151                 for (int i = 0; i < size; i++){
152                     if(!empties[i]){
153                         res=medians[i];
154                         break;
155                     }
156                 }
157                 break;
158             case 2:
159                 int max=0;
160                 for (int i = 0; i < size; i++){
161                     if(medians[i]>max)
162                         max=medians[i];
163                 }
164                 for (int i = 0; i < size; i++){
165                     if(empties[i]){
166                         medians[i]=max+1;
167                     }
168                 }
169                 serial_sort(&medians, size);
170                 res=medians[(size-all_empty)/2];
171                 break;
172             case 3:
173                 int sum=0;
174                 for (int i = 0; i < size; i++){
175                     if(empties[i]){
176                         sum += medians[i];
177                     }
178                 }
179                 if(size-all_empty!=0)
180                     res=sum/(size-all_empty);
181                 break;
182         }
183         for (int i = 0; i < size; i++)
184             medians[i]=res;
185         free(empties);
186     }
187     // scattering pivot value
188     MPI_Scatter(medians, 1, MPI_INT, &res, 1, MPI_INT, root, *comm);
189     free(medians);
190     return res;
191 }
192
193 void merge_lists(int* lis1, int* lis2, int len1, int len2, int** arr_pointer){
194     int* merge;
195     if(len1+len2==0){
196         merge=NULL;
197     }
198     else{
199         int i=0; int j=0;
200         merge = (int*) malloc((len1+len2)*sizeof(int));
201         while (i<len1 && j<len2)
202         {
203             if(lis1[i]<lis2[j]){
204                 merge[i+j]=lis1[i];
205                 i++;
206             }
207             else{
208                 merge[i+j]=lis2[j];
209                 j++;
210             }
211         }
212         while(i<len1){
213             merge[i+j]=lis1[i];
214             i++;
215         }
216         while(j<len2){
217             merge[i+j]=lis2[j];
218             j++;
219         }
220     }
221 }
222
223

```

```

224         free(*arr_pointer);
225         *arr_pointer = merge;
226     }
227
228     int main(int argc, char **argv) {
229         if (4 != argc) {
230             printf("Usage: qsort input_file output_file pivot_strategy (1-3)\n");
231             return 1;
232         }
233         char *input_name = argv[1];
234         char *output_name = argv[2];
235         int pivot_strategy = atoi(argv[3]);
236
237         int *local_data=NULL;
238         int *input=NULL;
239         int* output=NULL;
240
241         double run_time;
242         double max_runtime;
243         int size; int log2size;
244         int rank;
245         int n; int length; int length_shift=0;
246         int median, pivot;
247         int goal_idx, send_size, cnt, rec_size;
248         int idx_send, idx_use;
249         int* buffer=NULL;
250         int* sendc=NULL; int* recc=NULL; int* displ=NULL;
251         int i;
252         int is_empty=0;
253
254
255
256         MPI_Request send_req1;
257         MPI_Request send_req2;
258
259
260         //MPI init
261         MPI_Init(&argc, &argv);
262
263         MPI_Comm_size(MPLCOMM_WORLD, &size);
264         MPI_Comm_rank(MPLCOMM_WORLD, &rank);
265
266
267         if(!is_power_of_two(size))//checking if number of processor really is a power of 2
268             {printf("The number of processors shouldd be a power of 2!\n"); return 0;}
269
270         log2size = log_2(size);
271         MPI_Comm comm[log2size];
272         int ranks[log2size];
273         int sizes[log2size];
274         int color;
275
276         //Setting up the data
277         if(rank==root){
278             // Read input file
279             if (0 > (n = read_input(input_name, &input))) {
280                 return 2;
281             }
282         }
283     }
284
285         // Start timer
286         MPI_Barrier(MPLCOMM_WORLD);
287         double start = MPI_Wtime();
288
289
290         //ALG step 1 - Distributing data as equal as possible
291         MPI_Bcast(&n, 1, MPI_INT, root, MPLCOMM_WORLD);
292         length=n/size;
293
294         if(n<size){
295             if(rank==root){
296                 double start = MPI_Wtime();
297                 serial_sort(&input,n);
298                 run_time=MPI_Wtime()-start;
299                 output=input;
300

```

```

301     }
302 }
303 else{
304     if(n%size==0){ // Distribution if it is possible equally
305         local_data=(int*)malloc(length*sizeof(int));
306         MPI_Scatter(input, length, MPI_INT, local_data, length, MPI_INT, root,
307                     MPLCOMM_WORLD);
308     }
309     else{ // Distribution if it is impossible equally
310         length++;
311         local_data=(int*)malloc(length*sizeof(int));
312         if(rank >= n%size)
313             length_shift=-1;
314
315         sendc=(int*)malloc(size*sizeof(int));
316         displ=(int*)malloc(size*sizeof(int));
317         for (int i = 0; i < size; i++)
318         {
319             if (i < n%size)
320             {sendc[i]=length;
321              if(i!=0)
322                  displ[i]=displ[i-1]+length;
323              else
324                  displ[i]=0;}
325             else
326             {sendc[i]=length-1;
327              if (i==n%size)
328                  displ[i]=displ[i-1]+length;
329              else
330                  displ[i]=displ[i-1]+(length-1);}
331         }
332         MPI_Scatterv(input, sendc, displ, MPI_INT, local_data, length, MPI_INT, root,
333                     MPLCOMM_WORLD);
334         length=length+length_shift;
335         free(sendc);free(displ);
336     }
337     if(rank==root)
338         free(input);
339
340     // ALG 2 - Sort the data locally for each process
341     serial_sort(&local_data, length);
342
343     //ALG 3 - Perform global sort (Note even if it a recursive algorithms it is implemented
344     //withouth recursion as the depth is fix)
345     MPI_Comm_dup(MPLCOMM_WORLD,&(comm[0]));
346     for (i = 0; i < log2size; i++)
347     {
348         // Setting up the splitting of the communication
349         MPI_Comm_size(comm[i], &(sizes[i]));
350         MPI_Comm_rank(comm[i], &(ranks[i]));
351         color=ranks[i]/(sizes[i]/2);
352         MPI_Comm_split(comm[i], color, ranks[i], &(comm[i+1]));
353
354         //ALG 3.1 get pivot element with
355         if(length==0){
356             is_empty=1;
357             median=0;}
358         else{
359             is_empty=0;
360             median=local_data[(length)/2];
361             pivot=get_pivot(median, is_empty, &comm[i], ranks[i], sizes[i],
362                             pivot_strategy);
363
364         //ALG 3.2 divide data according to pivot
365         //finding splitting points
366         for (cnt = 0; cnt < length; cnt++){
367             if(local_data[cnt]>pivot){
368                 break;
369             }
370         }
371
372         //setting up indexes to exchange data
373         if(ranks[i] < sizes[i]/2){
374             goal_idx=ranks[i]+sizes[i]/2;
375             send_size=length-cnt;

```

```

374         idx_send=cnt;
375         idx_use=0;
376     }
377     else{
378         goal_idx=ranks[i]-sizes[i]/2;
379         send_size=cnt;
380         idx_send=0;
381         idx_use=cnt;
382     }
383
384     //ALG 3.3 Communicate data in the given pair
385     //Note: it is possible that some list is empty in this process, which
           needs some care
386     MPI_Isend(&send_size, 1, MPI_INT, goal_idx, 0, comm[i], &send_req1);
387     MPI_Recv(&rec_size, 1, MPI_INT, goal_idx, 0, comm[i], MPI_STATUS_IGNORE)
           ;
388     MPI_Wait(&send_req1, MPI_STATUS_IGNORE);
389
390     if (send_size!=0)
391         MPI_Isend(&(local_data[idx_send]), send_size, MPI_INT, goal_idx,
392                 0, comm[i], &send_req2);
393
394     if (rec_size!=0){
395         buffer=(int*) malloc(rec_size*sizeof(int));
396         MPI_Recv(buffer, rec_size, MPI_INT, goal_idx, 0, comm[i],
397                 MPI_STATUS_IGNORE);
398     }
399     if (send_size!=0)
400         MPI_Wait(&send_req2, MPI_STATUS_IGNORE);
401
402     //ALG 3.4 merge the local and received list - modify local length
           merge_lists(&(local_data[idx_use]), buffer, length-send_size, rec_size,
           &local_data);
403     length=length-send_size+rec_size;
404     free(buffer);
405     buffer=NULL;
406 }
407
408 //ALG 5 - Gather sorted data in the root
409 recc=(int*) calloc(size, sizeof(int));
410 displ=(int*) malloc(size*sizeof(int));
411
412 //Get the length of the arrays in the processes
413 MPI_Gather(&length, 1, MPI_INT, recc, 1, MPI_INT, root, MPLCOMM_WORLD);
414 displ[0]=0;
415 for (int i = 1; i < size; i++){
416     displ[i]=displ[i-1]+recc[i-1];
417 }
418
419 //Gather the final ordered list
420 if(rank==root)
421     output=(int*) malloc(n*sizeof(int));
422 MPI_Gatherv(local_data, length, MPI_INT, output, recc, displ, MPI_INT, root,
           MPLCOMM_WORLD);
423
424 //Stop clock
425
426 run_time = MPI_Wtime()-start;
427 //get the measured runtime
428 MPI_Reduce(&run_time, &max_runtime, 1, MPI_DOUBLE, MPI_MAX, root, MPLCOMM_WORLD
           );
429
430 // Clean up
431 free(local_data);
432 free(recc);
433 free(displ);
434 }
435
436 // Write results from root
437 if(rank==root){
438     printf("%lf\n", max_runtime);
439
440     #if save
441     if (0 != write_output(output_name, output, n)) {
442         return 2;
443     }

```



```

444         free(output);
445     #endif
446 }
447
448     if(n>=size)
449     for (int i = 0; i < log2size; i++){
450         MPI_Comm_free(&comm[i]);
451     }
452
453     MPI_Finalize();
454     return 0;
455 }

```

Makefile

```
1 #####
2 # Makefile for assignment 2, Parallel and Distributed Computing 2023.
3 #####
4
5 CC = mpicc
6 CFLAGS = -std=c99 -g -O3
7 LIBS = -lm
8
9 BIN = quicksort
10
11 all: $(BIN)
12
13 quicksort: quicksort.c
14             $(CC) $(CFLAGS) -o $@ $< $(LIBS)
15
16 clean:
17             $(RM) $(BIN)
```